

# Understanding, Detecting and Localizing Partial Failures in Large System Software<sup>1</sup>

October 16, 2020

---

<sup>1</sup>Chang Lou, Peng Huang, and Scott Smith. “Understanding, Detecting and Localizing Partial Failures in Large System Software”. In: *17th {USENIX} Symposium on Networked Systems Design and Implementation ({NSDI} 20)*. 2020, pp. 559–574.

# Overview

Problem definition

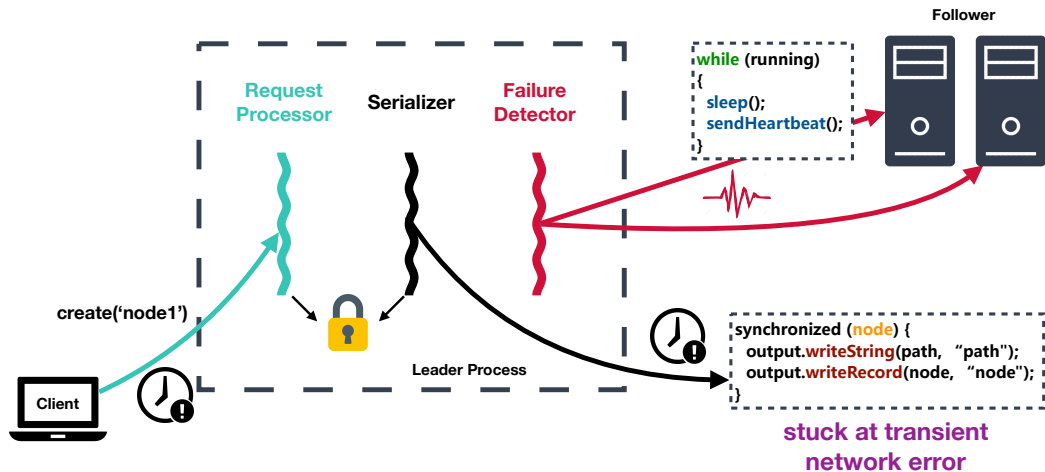
Case Study  
Findings

Motivation

Proposed Design  
Ideas  
Implementation

# What is a Partial Failure?

## An Example

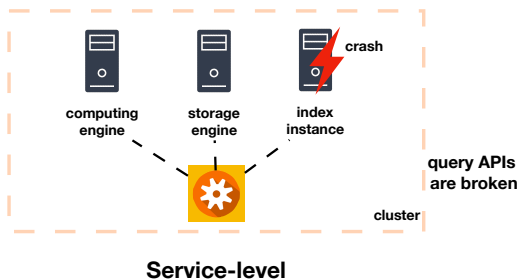
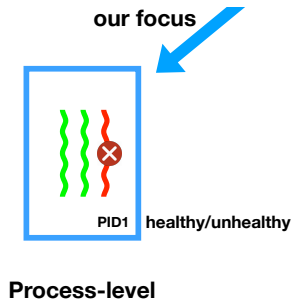


# What is a Partial Failure?

## Definition

A partial failure is, in a process  $\pi$  to be when a fault **does not** crash  $\pi$  but causes safety or liveness violation or severe slowness for some functionality  $R_f \subsetneq R$

**Scope:** In this paper, we will specify the partial failure at the **process** granularity instead of **service**.



# Study methodology

## 100 partial failure cases from five large, widely-used software systems

- ▶ Crawl all bug tickets tagged with critical priorities in the official bug trackers
- ▶ Filter tickets from testing and randomly sample the remaining failures tickets.

Interestingly, **54%** of them occur in the most recent **three** years' software releases  
(*average lifespan of all systems is 9 years*)


Software	Language	Cases	Versions	Date Range
ZooKeeper	Java	20	17 (3.2.1–3.5.3)	12/01/2009–08/28/2018
Cassandra	Java	20	19 (0.7.4–3.0.13)	04/22/2011–08/31/2017
HDFS	Java	20	14 (0.20.1–3.1.0)	10/29/2009–08/06/2018
Apache	C	20	16 (2.0.40–2.4.29)	08/02/2002–03/20/2018
Mesos	C++	20	11 (0.11.0–1.7.0)	04/08/2013–12/28/2018

# Finding 1: Root Causes are Diverse

Root cause distribution

**No single uniformed or dominating root cause<sup>2</sup>**

Top three (total 48%) root cause types are **uncaught errors**, **indefinite blocking**, and **buggy error handling**

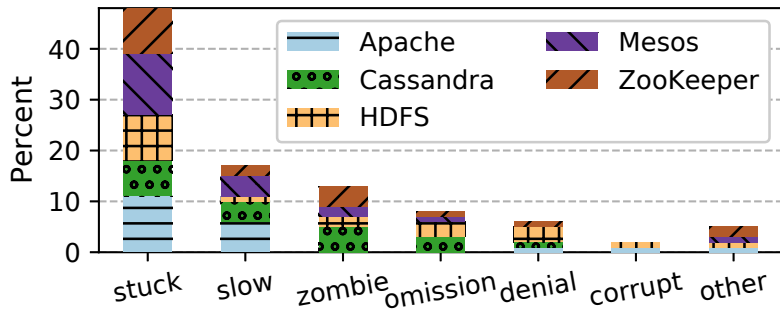


<sup>2</sup>UE: uncaught error; IB: indefinite blocking; EH: buggy error handling; DD: deadlock; PB: performance bug; LE: logic error; IL: infinite loop; RL: resource leak.

## Finding 2: Nearly Half Cases Cause **Stuck** Issues

### Consequence

Nearly half (**48%**) of the partial failures cause some functionality to be **stuck**.



**17%** of the partial failures cause certain operations to take a long time to complete.  
(i.e. **slow**)

## Other Findings: Partial Failures are Hard to Detect

**15%** of the partial failures are silent

Including data loss, corruption, inconsistency, and wrong results

**Most** cases are triggered by unique production workload or environment

**71%** of the partial failures are triggered by some **specific environment condition**, or **special input** in the **production**.

**Debugging time is long**

The median diagnosis time is 6 days and 5 hours

**The majority (68%) of the failures are “sticky”**

The process will not recover from the faults by itself. The faulty process needs to be restarted or repaired to function again.



# Motivation

So how to detect and localize a partial failure in a big software?

What if we simply apply static or dynamic analysis?

## Static Analysis?

- ▶ no unique production env/workload
- ▶ unable to detect run-time problem

## Dynamic Analysis?

- ▶ existing detectors are too shallow
- ▶ unable to localize failures

Ask developers to manually add defensive checks?


## Manual vs generated checkers

**Systematically** generated checkers to ease developers' burden


- ▶ challenge: difficult to automate for all cases
- ▶ opportunity: most of partial failures do not rely on deep semantic understanding to detect, such checkers can potentially be automatically constructed

# Intersection Principle

Construct customized checkers that **intersect** with the execution of a monitored process:



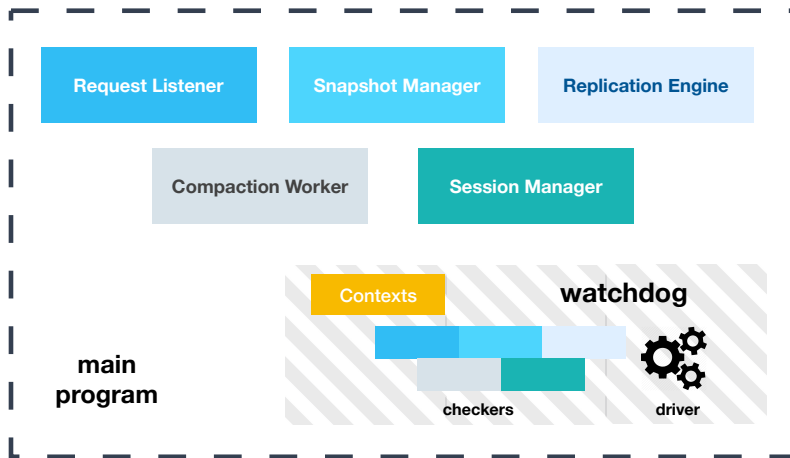
**existing approach**



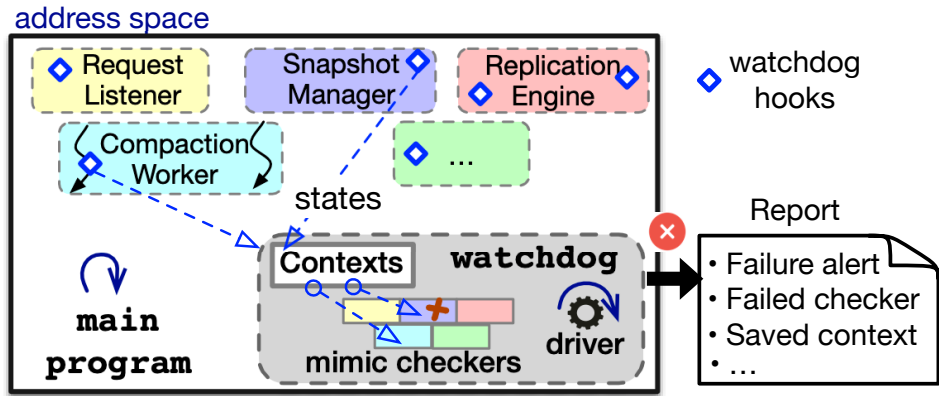
**our approach**

# Intrinsic watchdog: Runtime

An intrinsic watchdog is a dedicated monitoring extension for a process

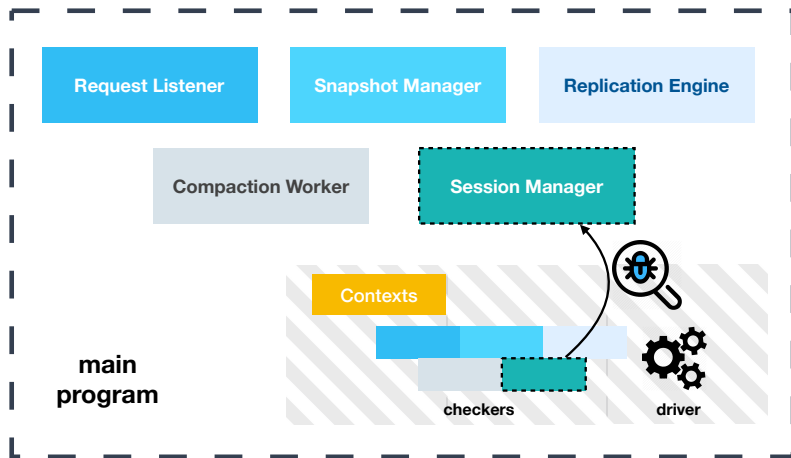


# Intrinsic watchdog: How it works?



# Characteristic I: Customized

- ▶ Regularly executes a set of checkers tailored to different modules
- ▶ Selects some representative operations from each module

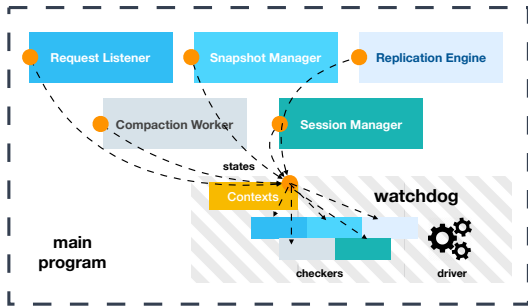


# Characteristic II: Stateful

To synchronized states, introduce

## Context

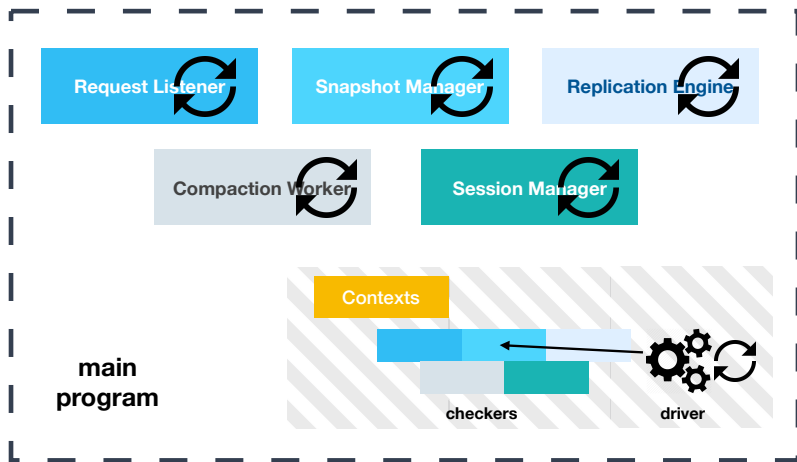
- ▶ bound to each checker
- ▶ holds all the arguments needed for the checker execution
- ▶ synchronized with the program state through hooks in the main program
- ▶ update with current state when hooks reached



**Note:** The watchdog driver will not execute a checker unless its context is ready.

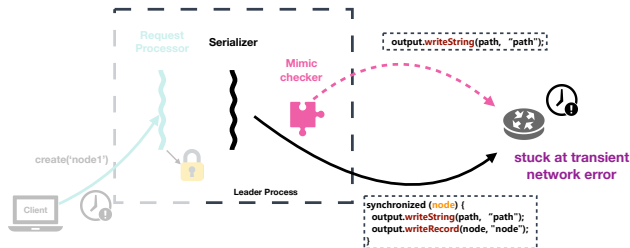
# Characteristic III: Concurrent

Run watchdog **concurrently** with the main program instead of **in-place** checking with **inserted** checkers



# Core Idea: Mimic Checking

Imitates some representative operations



**Exmample:** Perform a similar operation (snapshot) and also get stuck at the same location

## Accuracy

- ▶ exercises code logic similar to the main program
- ▶ share execution environment in runtime
- ▶ increases coverage of checking targets
- ▶ can pinpoint the faulty module and failing instruction



# Implementation: OmegaGen

## Tool Overview


- ▶ a prototype that systematically generates mimic-type watchdogs for system softwares
- ▶ in Java with 8,100 SLOC, using Soot analysis framework
- ▶ **core technique:** program reduction

```
1 $ ./omegagen -jar zookeeper-3.4.6.jar -m zookeeper.manifest
2 analyzing..
3 generating..
4 instrumenting..
5 repackaging..
6 done. Total 1min 6s.
7 $ ls output/
8 zookeeper-3.4.6-with-wd.jar
```

# What is Program Reduction?

## Definition

Given a program  $P$ , create a watchdog  $W$  that can detect partial failures in  $P$  without imposing on  $P$ 's execution.



**Reduce:** We need not put everything into checkers, because a lot of operations are logically **deterministic** and should be checked before production. Only some of them are more **vulnerable** in the production environment.

# Program Reduction

For the source code of a given program, the process will go through five steps:

1. locate long-running regions
2. reduce the program
3. locate vulnerable operations
4. encapsulate watchdog checkers
5. insert watchdog hooks

## Step 1: Locate Long-running Regions

Identifies potentially long-running loops in the function body

e.g. `while(true),while(flag)`

However, an identified **long-running** loop may turn out to be **short-lived** in an actual run.

### **predicate**-based algorithm

a runtime property associated with a method which tracks whether a call site of this method is in fact reached

- ▶ insert a hook before the loop → sets its predicate
- ▶ insert a hook after the loop → unset its predicate
- ▶ pass caller's predicate set to callees

Runtime:

- ▶ activate or activates or deactivates the associated watchdog based on assigned predicate

# Step 1: Locate Long-running Regions

An example

**initialization  
stage**

```
public class SyncRequestProcessor {  
    public void run() {  
        int logCount = 0;
```

```
        setRandRoll(r.nextInt(snapCount/2));
```

**long-running  
stage**

```
        while (running) {  
            ...  
            if (logCount > (snapCount / 2 ))  
                zks.takeSnapshot();  
        }
```

**cleanup  
stage**

```
        LOG.info("SyncRequestProcessor exited!");  
    }
```



## Step 2: Reduce the Program

Recursively analyze each function to find out **vulnerable** operations (in the next step)

```
public class SyncRequestProcessor {  
    public static void serializeSnapshot(DataTree dt, ...) {  
  
        ...  
        dt.serialize(oa, "tree");  
    }  
}  
  
public class DataTree{  
    public void serialize(OutputArchive oa, String tag) {  
        scout = 0;  
        serializeNode(oa, new StringBuilder(""));  
        ...  
    }  
}
```

The diagram illustrates the process of reducing the program by identifying recursive calls. Two code snippets are highlighted with red boxes and labeled "keep reducing". A red arrow points from the first box to the second box, indicating a recursive call.

## Step 3: Locate Vulnerable Operations

Looks for potentially **vulnerable** operations in the control flow of those long-running methods.

- ▶ Heuristics (default):
  - ▶ synchronisation
  - ▶ resource allocation
  - ▶ event polling
  - ▶ async waiting
  - ▶ invocations using external arguments
  - ▶ file or network I/O
  - ▶ complex while loop conditional
- ▶ Customize rule table in configuration
- ▶ Developers can also explicitly annotate an operation as `@vulnerable` in source codes

## Step 4: Encapsulate Watchdog Checkers

Construct reduced method for each **vulnerable** method in main program

```
public class SyncRequestProcessor$Checker {  
    public static void serializeNode_reduced(OutputArchive arg0, DataNode arg1) {  
        try{  
            arg0.writeRecord(arg1, "node");  
        } catch (Throwable ex) {  
            ...  
        }  
    }  
    public static Status checkTargetFunction0() {  
        ...  
        Context ctx = ContextFactory.serializeNode_reduced_context();  
        if (ctx.status == READY) {  
            OutputArchive arg0 = ctx.args_getter(0);  
            DataNode arg1 = ctx.args_getter(1);  
            executor.runAsyncWithTimeout(serializeSnapshot_reduced(arg0, arg1), TIMEOUT);  
        }  
        else  
            LOG.debug("checker context not ready");  
        ...  
    }  
}
```

extracted vulnerable  
operations



## Step 5: Insert Watchdog Hooks

To capture the real state of the main program in runtime and pass it to the checker

```
void serializeNode(OutputArchive oa, StringBuilder path) throws IOException {  
    String pathString = path.toString();  
    DataNode node = getNode(pathString);  
  
    String children[] = null;  
    synchronized (node) {  
        oa.writeRecord(node, "node");  
        Set<String> childs = node.getChildren();  
        if (childs != null)  
            children = childs.toArray(new String[childs.size()]);  
    }  
    path.append('/');  
    int off = path.length();  
    ...  
}
```

+ ContextFactory.serializeNode\_context\_setter(oa, node);

insert context hook before vulnerable operation

# An overview example

```
1 public class SyncRequestProcessor {
2     public void run() {
3         while (running) { 1 identify long-running region
4             if (logCount > (snapCount / 2))
5                 zks.takeSnapshot();
6             ... 3 reduce
7         }
8     }
9 }
10 public class DataTree { 3 reduce
11     public void serializeNode(OutputArchive oa, ...) {
12         ...
13         String children[] = null;
14         synchronized (node) { 2 locate vulnerable operations
15             scount++;
16             oa.writeRecord(node, "node");
17             children = node.getChildren();
18         }
19         ...
20     } + ContextManger.serializeNode_reduced
21 } _args_setter(oa, node);
4 insert context hooks
```

(a) A module in main program

```
1 public class SyncRequestProcessor$Checker {
2     public static void serializeNode_reduced(
3         OutputArchive arg0, DataNode arg1) {
4         arg0.writeRecord(arg1, "node");
5     }
6     public static void serializeNode_invoke() {
7         Context ctx = ContextManger. 4 generate
8         serializeNode_reduced_context(); context
9         if (ctx.status == READY) { factory
10             OutputArchive arg0 = ctx.args_getter(0);
11             DataNode arg1 = ctx.args_getter(1);
12             serializeNode_reduced(arg0, arg1);
13         }
14     }
15     public static void takeSnapshot_reduced() {
16         serializeList_invoke();
17         serializeNode_invoke();
18     }
19     public static Status checkTargetFunction0() {
20         ... 5 add fault signal checks
21         takeSnapshot_reduced();
22     }
23 }
```

(b) Generated checker